

Privacy in Statistics and Machine Learning
In-class Exercises for Lecture 16 (Zero Sum Games, Part I)
April 3, 2025

Spring 2025

Adam Smith (based on materials developed with Jonathan Ullman)

Problems with marked with an asterisk () are more challenging or open-ended.*

1. Consider a two-player zero-sum game described by a payoff matrix $M \in \mathbb{R}^{|\mathcal{R}| \times |\mathcal{C}|}$ and let $(\mathbf{r}, \mathbf{c})$ be a pair of equilibrium strategies. The support of a strategy is the set of actions with non-zero probability, so

$$\text{supp}(\mathbf{r}) = \{i : r_i > 0\}$$

and likewise for $\text{supp}(\mathbf{c})$. Prove that every i in the support of $\mathbf{r}$ is a best-response to $\mathbf{c}$. That is

$$\forall i \in \text{supp}(\mathbf{r}) \quad \mathbb{E}_{j \sim \mathbf{c}} (M_{i,j}) = \max_{i' \in \mathcal{R}} \mathbb{E}_{j \sim \mathbf{c}} (M_{i',j})$$

Note that the analogous statement (with min in place of max) will be true for all actions in the support of $\mathbf{c}$ by symmetry, but don't spend time proving it separately.

2. Consider the two-player zero-sum game with two actions for each player described by the payoff matrix

$$\begin{bmatrix} +2 & -1 \\ -2 & +3 \end{bmatrix} \quad (1)$$

Compute a pair of equilibrium strategies $(\mathbf{r}, \mathbf{c})$ for this game. [**Hint:** How does the property you proved in Question 1 help you find the equilibrium?]

3. The minmax theorem is the basis of one of the most widely used approaches for proving *lower bounds* on the performance of algorithms. Suppose we want to design an algorithm A that takes an input x from some finite set X and computes something about x . We have some real-valued *cost function* $\text{cost}(x, A)$ describing the cost of running A on input x . (This cost could be running time, space usage, query complexity, some measure of "error". It doesn't really make a difference.) For a given algorithm A , the *worst-case cost* of A is

$$\max_{\text{inputs } x} \text{cost}(x, A)$$

A randomized algorithm R can always be viewed as a distribution over deterministic algorithms. In that case, we can consider the *worst-case expected cost* of R as

$$\max_{\text{inputs } x} \mathbb{E}_{A \sim R} (\text{cost})(x, A)$$

- (a) (Yao's minimax principle, part 1) Show that if there exists a distribution $\mathbf{x}$ on inputs such that for every algorithm A

$$\mathbb{E}_{x \sim \mathbf{x}} (\text{cost}(x, A)) \geq T$$

then for every randomized algorithm A , the worst-case expected cost is at least T .

- (b) (Yao's minimax principle, part 2) Show that if every randomized algorithm R has worst-case expected cost at least T , then there exists a distribution on inputs $\mathbf{x}$ such that for every algorithm A

$$\mathbb{E}_{\mathbf{x} \sim \mathcal{D}} (\text{cost}(\mathbf{x}, A)) \geq T$$

In other words, if every (randomized) algorithm has to have high cost on some input, then there is a single distribution on inputs such that every (randomized) algorithm has to have high expected cost on that distribution.

Note: You can assume that the set of algorithms is finite. For example, they could be defined by Boolean circuits of some finite size.

[*Hint:* Consider a game in which the row player chooses an algorithm and the column player chooses an input.]

4. In a previous lecture we showed that MWEM is (ϵ, δ) -differentially private and can answer a set of k queries on a dataset in $\mathcal{U}^n$ with error at most $\leq \alpha$ on every query (with high probability), provided that

$$n \gtrsim \frac{(\log |\mathcal{U}|)^{1/2} (\log \frac{1}{\delta})^{1/2} (\log k)}{\epsilon \alpha^2}$$

Modify the analysis of the algorithm to ensure $(\epsilon, 0)$ -differential privacy? Prove a similar guarantee to above, showing that the algorithm is accurate provided that

$$n \gtrsim \frac{(\log |\mathcal{U}|)^a (\log k)^b}{\epsilon^c \alpha^d}$$

for some constants a, b, c, d . What parts of the algorithm and its analysis have to change?